

Technical Research Report - Sensing System

Graduation internship

Casey Stijlaart

Version 0.6, 08-09-2026

Table of Contents

- 1. Document history 1
- 2. Non-Camera Person Position & Height Sensing Techniques 2
 - 2.1. Introduction 2
 - 2.2. Research Approach 3
 - 2.3. Evaluation Criteria 4
 - 2.4. Candidate Sensor Technologies 5
 - 2.5. Module Selection 14
 - 2.6. Conclusion 15
- 3. End-to-End Latency Requirements for Natural Virtual Human Reaction 16
 - 3.1. Introduction 16
 - 3.2. Research Approach 17
 - 3.3. Perceptual Latency Thresholds 17
 - 3.4. Latency Budget 18
 - 3.5. Requirements and Verification 21
 - 3.6. Conclusion 22
- 4. Embedded Compute Platform Selection 23
 - 4.1. Introduction 23
 - 4.2. Research Approach 24
 - 4.3. Evaluation Criteria 25
 - 4.4. Candidate Embedded Platforms 26
 - 4.5. Conclusion 34
- 5. Bibliography 35

1. Document history

Table 1. Document History

Version	Date	Status	Author	Description	Remarks
0.1	03-09-2026	Draft Version	Casey Stijlaart	Draft Version	Initial document layout, non-camera sensing techniques chapter
0.2	04-09-2026	Draft Version	Casey Stijlaart	Draft Version	draft version of sensing techniques
0.3	07-09-2026	Draft Version	Casey Stijlaart	Draft Version	Finalisation of non-camera sensing techniques
0.4	07-09-2026	Draft Version	Casey Stijlaart	Draft Version	Initial draft of end-to-end latency chapter
0.5	08-09-2026	Draft Version	Casey Stijlaart	Draft Version	Finalised end-to-end latency
0.6	08-09-2026	Draft Version	Casey Stijlaart	Draft Version	Initial draft of embedded compute platform chapter

2. Non-Camera Person Position & Height Sensing Techniques

2.1. Introduction

This chapter addresses the first technical research question defined in the Project Plan (Chapter 3.6, Research Questions), referred to as RQ1: "Which combination of non-camera sensors provides sufficient 3D spatial resolution to distinguish adult and child height categories within a room-scale interactive space?" Two further sub-questions are referenced by number: RQ2 (robust, low-latency position estimates under real-world noise) and RQ4 (extensibility to multi-person scenarios); RQ3 (on-device ML for height classification) is out of scope here.

Per the Project Plan (Chapter 3.3, Assignment), the target system estimates a person's real-time 3D position and height within the Holobox space without cameras, for the Virtual Human to act on. This chapter surveys candidate technologies against the criteria they must satisfy, compares current hardware and cost, and recommends a sensor choice with a staged de-risking plan rather than a final hardware order, since a fully polished prototype is out of scope (Chapter 3.4, Scope and Preconditions). A general height estimate is treated as sufficient for now, and multi-person tracking is addressed only as an evaluation criterion, not a capability this stage builds.

2.2. Research Approach

This chapter follows the Library and Lab strategies of the DOT-framework identified for this research question in the Project Plan (Chapter 3.7, Research Methods): available product analysis, literature study, document analysis, and bench experimentation.

2.2.1. Library: Literature Study

Candidate Sensor Technologies gathers and verifies peer-reviewed literature for each non-camera sensing family against the position/height requirement and, separately, against the multi-person active-speaker problem, in two separate passes: one sits in the radar/thermal/RF position-and-pose-estimation literature, the other closer to acoustic and vibrometry-based speaker localisation, so a single combined search would have missed sources specific to either.

2.2.2. Lab: Available Product Analysis

Alongside the literature, the same chapter narrows the literature-backed technology families down to specific, currently sourceable hardware modules, checked against live distributor listings (Digi-Key, Mouser, Adafruit, and others) rather than assumed pricing, since an unverified cost estimate is not a usable basis for a hardware decision or the budget-approval process described in the Project Plan (Chapter 6.1, Finance and Risks). Literature backing and product availability are two different questions: a technique can be well-published without there being a specific, currently purchasable module that implements it at a reasonable price, or at all.

2.2.3. Lab: Bench Experimentation Strategy

The recommended path into bench experimentation commits to mmWave radar directly, managing implementation risk by splitting the build into capability-based phases rather than assuming the full pipeline works at once: position tracking first, then height estimation, then multi-person clustering, then speaker detection, each building on an already-working previous phase. This is deliberate: it isolates a failure to whichever capability is currently being added, rather than letting it surface somewhere in an untested, fully assembled system; the same mmWave hardware and UART/point-cloud/Unreal Engine bridge built in phase one is reused, not redone, in every phase after it; it follows the research-question sequence, validating position/height first before adding multi-person clustering and speaker detection; and it reflects feature priority, since those extensions improve the system without being required for the current single-person scope. A working fallback exists at each stage: if a later phase's capability underperforms, the earlier phase's version remains a functioning system rather than a dead end, mirroring the fallback the project's own risk register anticipates for both sensor selection and on-device ML performance.

2.3. Evaluation Criteria

Before diving into available sensor modules, it is of great importance to understand the requirements they need to meet. Thus the comparison can be set out against explicit, justified requirements rather than an implicit sense of what "good" looks like. Each criterion below is traced to either a technical sub-question described in the Project Plan (Chapter 3.6, Research Questions), a key feature named in the Project Plan (Chapter 3.3, Assignment), or a project constraint from the Project Plan (Chapter 6.1, Finance and Risks), rather than to a generic notion of sensor quality.

Table 2. Evaluation Criteria

Criterion	Description
Height-estimation capability	Must produce an actual height reading (alone or fused), not just (x, y) position, per RQ1's spatial-resolution requirement; a general estimate suffices for now.
Position accuracy and coverage	Position error and range must be usable at the Holobox's actual room-scale dimensions, since RQ1 covers (x, y) as well as height and a sensor that can't cover the space isn't usable regardless of height capability.
Real-time embedded feasibility	Interactive latency on an embedded platform (targeting <100 ms, hard ceiling ~1 s), as RQ2 requires low-latency embedded estimates; thresholds follow Nielsen's (1993) response-time limits, building on Miller (1968).
Multi-person extensibility	Shouldn't fundamentally break if a second person enters the space later, since RQ4 and the Assignment both name multi-person as a future-facing requirement.
Data minimisation / privacy	Output can't identify an individual and captures only what's needed for position and height, required by the ethical/legal sub-questions given that visitors are children aged 8-12 (AVG/GDPR).
Literature and community backing	Prior published work must show the specific capability required, not just the technology in general, reducing the risk of committing project time to proving something never demonstrated at all.
Hardware cost	Cost of the sensor and any paired second sensor, checked against sourceable listings; a real constraint since no project budget has yet been agreed with MindLabs.

These criteria are not weighted equally, and are not applied as a hard pass/fail filter. Height-estimation capability and data minimisation are treated as closer to hard requirements, since they follow directly from RQ1 and the ethical/legal sub-questions; the remaining criteria (sensor combination, cost, power/footprint) are weighed against each other in the chapter that follows, since a technology that is merely more expensive or harder to pair with a second sensor is not disqualified the way one with no height precedent or an identifying output effectively is.

2.4. Candidate Sensor Technologies

This chapter surveys non-camera sensing technologies capable of device-free person detection, localisation, and (where the literature supports it) height estimation, evaluated against the technical sub-questions described in the Project Plan (Chapter 3.6, Research Questions): primarily sufficient 3D spatial resolution for height, embedded-platform feasibility, and the data-minimisation/privacy constraints that follow from the primary visitors being children aged 8-12. A general height estimate is sufficient for now; distinguishing adult and child height categories is a possible later refinement, not a hard requirement.

Each technology is covered once, in two parts: what the literature has actually shown it can do, and, for the technically viable candidates, how that holds up against the criteria set out earlier, specific hardware, cost, and pros/cons, closing in a comparison summary, criteria matrix, and recommendation. Ngamakeur et al. (2020) frame this same comparison at a broader scale, structuring device-free sensing into a pipeline (Figure 1) used below to organise what each technology has actually been shown to do, rather than what it could in principle do.




Figure 1. Device-free Pipeline

2.4.1. Sensing Technologies

Where a technology has literature for both problems, its write-up splits into Position & Height Localisation and Speaker Identification subsections; several technologies are covered for localisation only, and Laser Doppler Vibrometry and the Microphone Array for speaker identification only. None of the speaker-identification sources have been tested in this project's specific context (children's voices, a room-scale interactive installation), so they remain a lead for further work rather than a settled result.

2.4.1.1. Time-of-Flight & LiDAR

2D LiDAR measures range by timing a laser pulse's reflection while scanning a single horizontal plane, producing an (x, y) point cloud at the unit's mounted height; well suited to in-plane tracking, but a single fixed-height plane does not by itself yield a height measurement (would need multiple units at different heights, a tilted scan, or a full 3D ToF/LiDAR unit).

Within its native (x, y) role, 2D LiDAR performs well: Bouazizi et al. (2024) fuse multiple units with a ConvLSTM model and show tolerance to furniture occlusion. For multi-person settings, Alam et al. (2021) document "crossover ambiguity" as 2D LiDAR's recurring multi-person failure mode when tracks pass close together, the concrete risk to plan around later. The representative module, the RPLIDAR A2M8, scans at 5-15 Hz (typically 10 Hz) per Slamtec's datasheet, the limit on full (x, y) frame availability regardless of its faster ~4000 samples/sec per-sample rate.

Pros

- Strong in-plane position accuracy and furniture-occlusion tolerance when multiple units are fused (Bouazizi et al.).
- Mature, widely available consumer/robotics hardware (RPLIDAR family) with broad community support.

Cons

- No height reading from a single scan plane; needs a second sensor or added mechanical complexity to get one.
- "Crossover ambiguity" is a documented multi-person failure mode when tracks pass close together (Alam et al.).
- 5-15 Hz scan rate limits how quickly a full (x, y) frame becomes available.

2.4.1.2. mmWave Radar

mmWave radar has an extensive body of research literature covering both position tracking and skeletal/joint-level pose estimation, and its point-cloud output is inherently non-identifying (no recognisable imagery), aligning with the sub-questions on minimising personally identifiable data from child visitors. Available product analysis narrows this to a 60 GHz point-cloud radar module built around the TI IWR6843 chip, over 24 GHz presence sensors (binary presence flag only, no point cloud) and undocumented generic modules lacking reference software. Several IWR6843 evaluation-board variants exist at different price points and field-of-view characteristics; **Module Selection** compares them and picks one to carry forward.

Position & Height Localisation

For position tracking, Huang et al. (2021) combine clutter removal, clustering, and Kalman filtering for real-time multi-person tracking on a Raspberry Pi 4, and Pegoraro and Rossi (2022) report position error down to roughly 7.6 cm. Pose-estimation literature is also relevant here: Zhu et al.'s ProbRadarM3F (2024) recovers 14 body keypoints from radar point clouds, including head and shoulder joints, backed by a camera-free radar dataset with joint ground truth (Su et al., 2024). This project's own planned approach is simpler, though: just take the top and bottom point of a person's point cloud and measure the gap (see End-to-End Latency Requirements, Compute Platform Implications). It's cheaper and faster, but unlike the keypoint model, it isn't backed by literature yet. The keypoint model stays the fallback if the simple version doesn't hold up. On latency: reported tracking papers give broadly similar tens-of-milliseconds frame periods on different radar chips (15-20 fps). TI doesn't publish a fixed maximum frame rate for the IWR6843 chip itself, the mmWave SDK exposes frame periodicity as an open, configurable value ([**mmwavesdkug**]), not a spec. TI's own indoor people-tracking reference design for this chip ([**tidue71d**]) runs at roughly 18 Hz (55ms), in the same range as the 10-20 Hz "typical demo" configurations noted elsewhere; see Evaluation Criteria for the thresholds this is checked against.

Speaker Identification

Whether the IWR6843 reports two people as distinct targets depends on angular and range resolution, not a blanket "multi-person" capability: roughly 1.7-2.9m of horizontal separation, or a ~1.7m height difference, is needed before angular resolution tells them apart. Range separation is far finer, ~4-5cm, from time-of-flight alone.

mmWave's own speaker-ID path is Doppler vibrometry: sensing larynx skin-surface displacement via a ~90-200Hz bandpass. Xu et al.'s WaveEar (2019) demonstrates this, but only at short range; no source shows it working room-wide at the scale this project needs, so it's a bonus capability here, not the primary mechanism. In this design it confirms whoever is close to the Holobox (~1-1.5m), doubling as a liveness/anti-spoof check, a played-back recording has no matching throat-vibration signature.

Room-wide speaker identification is delegated to a paired **Microphone Array** instead, steered by the radar's per-person tracked angles. For an illustrative 3×3m tracked zone with 2-3 people, ordinary spacing satisfies counting and height; at this scale two people already approach the 1.7-2.9m threshold at the zone's diagonal (~4.2m), so ambiguous tracks are reachable with fewer people than a larger space would allow. That zone size is a planning illustration, not a confirmed Holobox spec.

Pros

- Extensive literature base: real-time multi-person tracking and keypoint-level pose/height estimation.
- Demonstrated real-time operation on a Raspberry Pi 4.
- Point-cloud output is inherently non-identifying, fitting the project's data-minimisation requirements.
- Also supports close-range (~1-1.5m) speaker confirmation and liveness/anti-spoof checking via Doppler vibrometry, without additional hardware, alongside the room-wide microphone array.

Cons

- Sensor cost is €160-185 for the IWR6843AOPEVM variant (see **Module Selection**).
- Height estimation still needs extra work: either the simple top/bottom-point calculation or a keypoint-extraction model. Neither is a single, settled method yet.
- No project budget has yet been agreed with MindLabs, so this cost is a real, unresolved risk.
- Radar-only speaker identification is not demonstrated room-wide; a paired microphone array is needed for that (see below).

2.4.1.3. Wi-Fi CSI Sensing

Wi-Fi Channel State Information (CSI) sensing infers presence and movement from fluctuations in the wireless channel caused by a body moving through it, using only commodity Wi-Fi hardware, but the gathered literature is consistent that it trades away spatial resolution and stability: CSI is sensitive to environmental changes (furniture, multipath, other RF sources).

Position & Height Localisation

IndoTrack (2017) is the foundational device-free CSI tracking paper in this space; more recent work narrows toward multi-person settings, e.g. Shen et al.'s time-selective RNN for busy-room detection (2023). Reported update rates vary by an order of magnitude across otherwise comparable commodity setups, from 10 Hz (Shen et al., capped by their router) to a reported 100 Hz (IndoTrack, unconfirmed against the primary paper), an additional latency risk on top of the lack of height output.

Speaker Identification

Wi-Fi CSI reflections from mouth/lip movement during speech are detectable using only commodity Wi-Fi: Wang et al.'s WiHear (2014) uses directional antennas to sense multiple people's talk simultaneously via MIMO.

Pros

- Effectively free: reuses existing Wi-Fi infrastructure.
- Multi-person and multi-room literature exists (Time-Selective RNN).
- Device-free output with no recognisable imagery.
- The same commodity hardware also has a literature precedent for detecting mouth-movement reflections for speaker identification (WiHear).

Cons

- No published height-estimation precedent at all, the single biggest gap for this project's needs.
- Sensitive to environmental changes (furniture, multipath, other RF sources), the primary limit on stability.
- No demonstrated real-time embedded deployment in the sources reviewed.

2.4.1.4. Ultrasonic Ranging

Active ultrasonic ranging measures distance via time-difference-of-arrival (TDOA) or time-of-flight on emitted pulses, sidestepping the "room full of talking people" interference problem a passive-listening approach would face.

Segers et al. (2015) track up to four transmitters at 1-2cm accuracy on an FPGA using frequency-shift-keyed ultrasound. Segers et al. do not report an update rate or latency figure, so latency here is an open question rather than a documented one.

Pros

- Sidesteps the "room full of talking people" interference problem passive listening would face.
- Demonstrated real-time operation on embedded hardware (an FPGA, per Segers et al.).
- Reasonable literature base for ranging and positioning.

Cons

- No gathered source addresses height or elevation.
- Full coverage needs multiple synchronised nodes, reintroducing multi-sensor complexity.
- Custom-built hardware, roughly €200-400 (rough estimate, not a single purchasable SKU).

2.4.1.5. Infrared / Thermal

Low-resolution thermal/thermopile arrays measure a grid of per-pixel temperature readings: a person appears as a warm blob, and the same low pixel count that limits resolution and range doubles as a privacy strength, too coarse to resolve recognisable imagery.

Naser et al. (2021) extend thermal-array distance estimation to height using a neural network, the most directly relevant source, and should be read first. Shubha and Shastrimath (2022) show 99% occupancy accuracy on a Grid-EYE/AMG8833 array on a Raspberry Pi; Klir et al. (2025) extend the same sensor class to multi-person localisation. Naser et al.'s sensor samples at 0.5-64 Hz; the AMG8833 breakout identified for procurement is coarser, only 10 Hz or 1 Hz.

Hardware cost is around $2 \times \text{€}40\text{-}55$, confirmed at €49.60 (incl. VAT) via [Kiwi Electronics](#) (NL); the I2C interface is straightforward. That price is for the AMG8833's narrow $\sim 60^\circ$ FoV, enough only for a small room; room-scale coverage without mounting overhead, not an option for the Holobox, instead means the wider 110° MLX90640 variant, €75.01 (incl. VAT) at [Kiwi Electronics](#) (NL). Paired with a room-scale primary sensor ([Combined Sensor Approaches](#)), two narrow-FoV AMG8833 units ($\sim \text{€}100$ total) split coverage instead of one pricier wide-FoV part.

Pros

- Height estimation is demonstrated directly from sensor data, not just a proxy (Naser et al.).
- Low sensor-only cost (€40-55 narrow-FoV; €70-75 for the wide-FoV variant needed for room-scale coverage).
- Demonstrated real-time embedded operation on a Raspberry Pi.
- Low pixel count doubles as a privacy strength: output cannot resolve recognisable imagery.

Cons

- Effective range constrains it to smaller rooms unless the wider-FoV variant is used.
- Lower spatial resolution overall, which needs checking against the Holobox's actual dimensions.

2.4.1.6. Capacitive & Floor-based Sensing

Capacitive / Electric-Field Sensing

Two device-free variants exist: proximity-plate sensing (measuring distance to a body via a change in capacitance) and floor-embedded capacitive electrodes (requiring a person to stand on the sensed area). Bian et al.'s survey (2024) covers the design space broadly, and Faulkner et al.'s CapLoc (2020) is a floor-embedded mat evaluated in the EvAAL indoor-localisation competition. Range is the main constraint: plate sensors top out around 0.3-1m, so whole-room coverage needs a floor grid or many nodes, and the signal is inherently presence/proximity rather than spatial, so it is not naturally suited to a height reading at all.

Floor-based (Pressure) Sensing

Floor-based sensing (pressure, piezoresistive, or capacitive-floor mats) is inherently a 2D floor-plane measurement, but because each footstep is a spatially discrete event, it naturally disambiguates multiple people. Gatto et al. (2015) use a 64x64 pressure-switch array per tile, and Ibara et al. (2013) show gait-based individual identification from stride length and direction.

Pros

- Floor-based variants trivially solve "how many people, where."
- Solid literature base for both sub-variants.
- Inherently non-identifying (presence/proximity or pressure, not imagery).

Cons

- No height signal at all from either sub-variant.
- Capacitive sensing has short range, needing a dense grid for whole-room coverage.
- Floor-based sensing requires physical installation under the existing floor; cost can run into the thousands of euros.

2.4.1.7. Laser Doppler Vibrometry

Speaker identification only; not evaluated for position/height localisation.

A laser Doppler vibrometer (LDV) measures skin-surface vibration at a distance via Doppler shift, picking up throat vibration during speech even in noisy conditions; Avargel and Cohen (2011) combine an LDV sensor aimed at the larynx with a standard microphone to enhance speech quality under transient noise.

2.4.1.8. Microphone Array (Passive Acoustic)

Speaker identification only; not evaluated for position/height localisation.

Berghi and Jackson (2023) localise an active talker from multichannel audio alone using a CRNN, no camera or radar needed. This project doesn't need that open-ended version, though: the radar pipeline (mmWave Radar) already tracks each person's position.

Room-wide Speaker Identification (Primary)

Instead of blind direction-of-arrival estimation, the array steers a few listening beams at the radar's tracked angles and compares acoustic energy across them to attribute speech, considerably more reliable than a blind approach, but dependent on radar tracking already working. An ML refinement of that beam comparison, rather than the blind CRNN approach, is worth exploring too; not yet evaluated for feasibility here. This makes the microphone array, not the radar's own close-range Doppler vibrometry, the primary path to room-wide speaker identification, covering the room the radar's ~1-1.5m range can't reach.

Open Risk: Array Range vs. Tracked-Zone Coverage

Off-the-shelf arrays like the Sseed ReSpeaker are rated for ~3m. For the illustrative 3×3m tracked zone used elsewhere, an edge mount's worst case (~4.2m, the zone's diagonal) exceeds that; a central mount's (~2.1m) doesn't. Both the zone size and mounting position are undecided, so this stays open: a higher-gain array or a central mount are the two directions to check before hardware is committed to.

2.4.2. Combined Sensor Approaches

The individual technologies above are also weighed as sensor combinations, where pairing two of them addresses a gap neither has alone.

2.4.2.1. 2D LiDAR + Thermal Fusion

As with the other thermal-fusion options above, the thermal side here is two AMG8833 units (~€100 total), splitting room coverage between them rather than relying on one pricier wide-FoV MLX90640 or a single narrow-FoV unit that falls short on its own. The representative LiDAR module itself, the RPLIDAR A2M8, also now shows as discontinued/out of stock at major distributors (e.g. Sseed Studio), an added sourcing risk on top of the concerns below.

Pros

- 2D LiDAR contributes strong in-plane position accuracy and occlusion tolerance (Bouazizi et al.).
- Thermal contributes the same direct height-estimation precedent used elsewhere in this chapter (Naser et al.).

Cons

- Sensor-only cost of €350-390, driven up by needing two thermal units.
- Combining two independent position-sensing streams from scratch introduces synchronisation and fusion complexity flagged as a disproportionate implementation risk for this project's timeline, as described in the Project Plan (Chapter 4.3, Timeline).
- The representative RPLIDAR A2M8 module is discontinued at major distributors, an unresolved sourcing risk specific to this combination.
- Does not reuse an already-working single-sensor pipeline; both position-sensing streams are new to this combination.

2.4.3. Comparison Summary

Table 3. Candidate Sensor Comparison Summary

Technology	Representative Module	Approx. Hardware Cost (sensor only)	Height Data in Literature	Documented Latency / Update Rate
mmWave Radar	TI IWR6843AOPEVM (see Module Selection for the variant comparison)	€174,92-€248.23	Yes, keypoint-based	10-20Hz as base, but configurable to higher rates.
Infrared/Thermal Array	AMG8833 breakout, Kiwi Electronics (NL)	€48.39	Yes, direct	1 Hz or 10 Hz
2D LiDAR + Thermal (sensor fusion)	RPLIDAR A2M8 + 2x AMG8833, Kiwi Electronics (NL)	€385.69 + 2x €48.39	Partial, via fusion	5-15 Hz
Wi-Fi CSI Sensing	None. It reuses existing commodity Wi-Fi hardware	Approx. €0 (existing infrastructure)	None found	10-100 Hz
Ultrasonic Ranging	Custom transducer array (e.g. SPM0404UD5 MEMS elements, per Segers et al. (2015))	€200-400 (rough estimate for custom-built hardware, not a single purchasable SKU)	None found	Not reported
Capacitive & Floor-based Sensing	Custom-built capacitive/pressure floor mat. There is no off-the-shelf module	€80-2000+, depending on variant (rough estimate for custom-built hardware)	None found	Not reported

Figures above are sensor/module hardware only. All candidates other than Wi-Fi CSI also need an embedded compute platform to run the sensing pipeline; that cost is excluded here since it is a shared requirement across candidates rather than something that distinguishes one sensor choice from another. Latency figures are the sensor's own documented output rate, not a measured end-to-end pipeline latency; see the caveat already noted in Evaluation Criteria.

Table 4. Criteria Comparison Matrix

Technology	Height Est.	Embedded Real-Time	Multi-Person	Fusion-Ready	Privacy	Literature	Price Range (sensor only)
mmWave Radar	✓	✓	✓	✓	✓	✓	€160-250
Infrared/Thermal Array	✓	✓	✓	✓	✓	✓	€80-110
2D LiDAR + Thermal Fusion	✓	□	✓	✓	✓	✓	€350-390
Wi-Fi CSI Sensing	□	□	✓	□	✓	✓	~€0 (existing infrastructure)
Ultrasonic Ranging	□	✓	□	□	✓	✓	€200-400
Capacitive & Floor-based Sensing	□	□	✓	□	✓	✓	€80-2000+

2.4.4. Recommended Approach

On the criteria set out earlier, mmWave radar comes out ahead. It's the only approach with keypoint-level pose literature, and it has the strongest literature base of any candidate reviewed. The infrared/thermal array is cheaper and has a direct height-estimation precedent, but mmWave is still the stronger technical choice. This is a recommendation of the modality; the specific IWR6843 variant is addressed in **Module Selection**. Height itself is only a general estimate for now, adult/child categorisation is a possible later refinement, not something this recommendation locks in.

For the RQ4 speaker-identification problem, radar alone only reaches ~1-1.5m around the Holobox, via Doppler vibrometry, which doubles as a liveness/anti-spoof check. Room-wide speaker identification is instead handled by a paired **Microphone Array**, steered by the angles radar already tracks for each person, rather than solving open-ended source localisation on its own.

mmWave is by far the most expensive candidate reviewed, and no project budget has been agreed with MindLabs yet, so clearing that cost with the company mentor is a precondition before ordering hardware. Implementation risk is managed by building the pipeline one capability at a time: position, then height, then multi-person clustering, then speaker detection, each on an already-working previous phase. The microphone array gets added once radar tracking works, since beam-steering depends on it, and only once its own hardware-range risk (see **Microphone Array**) is resolved.

2.5. Module Selection

Candidate Sensor Technologies recommends mmWave radar as the sensing modality, built around the TI IWR6843 chip. The chip ships on several evaluation-board variants at different price points and field-of-view characteristics, given below as azimuth (angular coverage side-to-side) / elevation (angular coverage top-to-bottom, i.e. floor-to-head reach). This chapter compares those variants directly and picks one. The retrieved product listings are current as of 07-09-2026 and may vary when checking on a different date.

- **IWR6843ISK**: the standard long-range antenna plug-in module, €248,23 at current **Mouser NL** listings, with a 120° azimuth / a narrower 30° elevation field of view tuned for range over coverage, which works against a room-scale height-estimation use case rather than for it.
- **IWR6843ISK-ODS**: a wide-field-of-view (120° azimuth / 120° elevation) antenna variant, etched on a Rogers substrate, that TI markets specifically for overhead/ceiling-mounted people counting and tracking (~12 m range), €205,44 at current **Mouser**.
- **IWR6843AOPEVM/AWR6843AOPEVM**: an antenna-on-package (AoP) variant, with the antenna integrated directly into the chip package rather than etched onto a separate PCB, €174,92 at current **Mouser NL/AWR6843AOPEVM**. Matches the ISK-ODS's wide 120° azimuth / 120° elevation field of view in a smaller (~65×60mm).

Table 5. mmWave EVM Variant Comparison

Feature	IWR6843ISK	IWR6843ISK-ODS	IWR6843AOPEVM/AWR6843AOPEVM
Price (€)	248,23	205,44	174,92
Antenna	60-64GHz	60-64GHz	60-64GHz
Rx/Tx	4/3	4/3	4/3
Azimuth Angular Resolution (deg)	15	29	29
Azimuth FOV (deg)	120	120	120
Elevation Angular Resolution (deg)	58	29	29
Elevation FOV (deg)	30	120	120
Max Distance for child/adult (+/-m)	50/60	40/48	41/49

NOTE

The maximum distance for the TI modules is given by the TI **mmWave Sensing Estimator**, configured for the short-range default; it is therefore theoretical, not a measured figure.

2.5.1. Recommendation

Of the three variants in [Table 5](#), the IWR6843AOPEVM (equivalently, AWR6843AOPEVM) is recommended. Elevation, not azimuth or range, is the deciding factor. Elevation FOV and resolution determine how much of a person's vertical extent falls inside the beam, the precondition for the keypoint-based height estimate. Every variant already covers the room's short range with margin to spare. Azimuth isn't a deciding factor either. All three variants share roughly the same 29° figure anyway (bar IWR6843ISK's finer 15°), and the fine angular discrimination room-wide speaker identification needs is handled by a paired microphone array instead (Candidate Sensor Technologies, [Microphone Array](#)), not the radar's own azimuth resolution.

IWR6843ISK is ruled out: its 30° elevation FOV is by far the narrowest of the three, its elevation resolution the coarsest, and it is also the most expensive at €248,23. IWR6843ISK-ODS matches AOPEVM's elevation FOV and resolution exactly, but costs 17% more and its antenna is built for ceiling mounting, an assumption the Holobox has not committed to. AOPEVM reaches the same elevation performance in a smaller, mounting-agnostic form factor at the lowest price of the three, which is why it is carried forward into the phased mmWave implementation in the Conclusion.

2.6. Conclusion

Candidate Sensor Technologies and Module Selection recommend mmWave radar, specifically the IWR6843AOPEVM: the only approach with keypoint-level pose literature for height, non-identifying output, and the strongest literature base reviewed, at the lowest price of the elevation-capable variants compared.

Implementation risk is managed by building the pipeline in phases, position, then height, then multi-person clustering, then speaker detection, each on an already-working previous phase; mmWave is the most expensive candidate reviewed, so mentor approval of that cost is a precondition. Sensor fusion (mmWave or 2D LiDAR with thermal) wasn't adopted as the starting point: mmWave alone already has a plausible path to height, no fusion combination is tested together in the literature, and thermal needs two units to match mmWave's coverage. Fusion stays a documented fallback, not a precondition for starting.

Speaker identification (RQ4) splits across two sensors. mmWave's own Doppler-vibrometry path only reaches ~1-1.5m around the Holobox, a bonus that doubles as a liveness check, not a room-wide solution. Room-wide speaker-ID is handled by a paired microphone array instead, steered by angles the radar already tracks (Candidate Sensor Technologies, [Microphone Array](#)), added once radar tracking works. Open risk: off-the-shelf arrays like the Sseed ReSpeaker are rated for ~3m. Whether that's enough depends on mounting position relative to the tracked zone (illustrative 3×3m example): an edge mount's worst case (~4.2m) exceeds that rating, a central mount's (~2.1m) doesn't. Worth a higher-gain array or committing to central mounting before hardware is ordered.

3. End-to-End Latency Requirements for Natural Virtual Human Reaction

3.1. Introduction

This chapter addresses the User/Human Interaction sub-question defined in the Project Plan (Chapter 3.6, Research Questions), referred to as RQ8: "What position and height estimation latency is required for the Virtual Human's reaction (e.g. kneeling down for a child) to be perceived by visitors as natural and immediate, rather than delayed or robotic?"

Per the Project Plan (Chapter 3.3, Assignment), the Virtual Human must react to a visitor's detected position and height so it reads as immediate and natural, not delayed or robotic, downstream of the sensing pipeline covered in Non-Camera Person Position & Height Sensing Techniques. Building or verifying that reaction isn't this project's job, though: Finished Products lists only the embedded sensing system as the prototype deliverable, the Virtual Human's own rendering is MindLabs' existing Holobox/Unreal Engine system. So this chapter sets two things: a full-pipeline target for RQ8 (sensing all the way to the visible reaction, derived from general HCI response-time literature and one telecom standard, since no literature specific to this reaction exists), and a narrower NFR that this project actually builds and can bench-test, sensing through to exposing the estimate to the pipeline. It breaks that budget down using the sensor already recommended there ([Section 2.4.1.2](#), [Section 2.5](#)). As with the sensing research, this chapter proposes a target and a verification plan, not a measured result. The bench and field measurement in Requirements and Verification hasn't run yet; that's the next step, not something this desk-research stage claims to have already confirmed.

3.2. Research Approach

This chapter follows the Library and Field strategies of the DOT-framework identified in the Project Plan (Chapter 3.7, Research Methods): literature study and document analysis to derive and budget the latency target.

3.2.1. Library: Literature Study

The "feels natural" requirement is grounded in general response-time literature rather than a number picked from intuition: Miller (1968), Nielsen (1993), and Shneiderman & Plaisant (2010) establish the standing 0.1s/1s/10s interactive-system response-time thresholds, and ITU-T G.114 (2003) gives a comparable, independently-derived threshold (150ms one-way) for real-time telecom traffic. A targeted search for literature on position/height-triggered avatar-reaction latency specifically, as opposed to general system responsiveness, did not surface a source that could be verified against its own claim; Perceptual Latency Thresholds is explicit about this gap and derives a project-specific target from the general literature instead of attributing a number to a source that does not actually make that claim.

3.2.2. Field: Document Analysis

Latency Budget decomposes the sensing-to-reaction pipeline into per-stage budgets (sensor acquisition, fusion, inference, rendering, optional transmission) by analysing existing documentation rather than running new experiments: the sensor acquisition figures reuse the recommended mmWave module's own datasheet and SDK configuration options (Section 2.4.1.2, Section 2.5), and the remaining stages are estimated from published figures for comparable embedded pipelines, since this project has not yet built and measured them.

3.3. Perceptual Latency Thresholds

3.3.1. General Interactive-System Thresholds

Miller (1968) establishes, and Nielsen (1993) and Shneiderman & Plaisant (2010) both restate, a three-tier response-time structure for interactive systems:

- **≤ 0.1s**: perceived as instantaneous; the user feels direct control over the system.
- **0.1s - 1s**: a noticeable but acceptable delay; the user's flow of thought stays uninterrupted.
- **1s - 10s**: noticeable enough that the system should give explicit feedback (e.g. a progress indicator), or the user's attention starts to drift.
- **> 10s**: the user disengages from the task entirely.

These thresholds describe general system responsiveness (a button press, a page load), not a reciprocal physical reaction between two agents specifically; the Holobox's Virtual Human reacting to a visitor's detected position and height (e.g. kneeling for a child) is closer to the latter, so the next section narrows the target rather than applying the general 1s "acceptable" threshold as-is.

3.3.2. A Comparable Real-Time Threshold: ITU-T G.114

ITU-T Recommendation G.114 (2003), the telecom standard for one-way transmission time in real-time voice/video traffic, independently arrives at a similar boundary from a different domain: below 150ms one-way delay is rated acceptable for most real-time applications, 150-400ms is acceptable but shows increasing impairment, and above 400ms is not acceptable. G.114 measures one-way, mouth-to-ear delay for telephony and video calls; it does not cover embodied position/height-reaction latency, and is cited here only as a second, independently-derived data point that a sub-150ms target sits inside a boundary already treated as "real-time acceptable" elsewhere, not as a source that makes a claim about virtual humans.

3.3.3. Project Target: Position/Height-Reaction Latency

Neither source above speaks to this project's specific reaction, position and height detection driving a Virtual Human behaviour change, so this project sets its own target rather than borrow one that doesn't fit: < 150ms end-to-end (sensing to visible reaction), with < 100ms as a stretch target. The reasoning: a reciprocal physical response, the Virtual Human kneeling for a detected child, sits inside Miller/Nielsen/Shneiderman & Plaisant's "acceptable" band (0.1-1s) only at its very low end. The visitor isn't issuing a command and waiting for a result the way those thresholds assume; they're perceiving a continuous, two-way physical interaction, closer to G.114's mouth-to-ear immediacy than to a page-load delay. 150ms is therefore adopted as the hard ceiling for the full pipeline, matching G.114's "acceptable" boundary, with 100ms as the stretch target matching the general literature's "instantaneous" threshold. This project only builds and can verify part of that pipeline, though, so Requirements and Verification turns this full-pipeline target into a narrower, actually-testable NFR for the part this project owns.

3.4. Latency Budget

3.4.1. Pipeline Breakdown

The pipeline this project actually builds and can verify is sensing → fusion → inference → output (the estimate exposed to the Unreal Engine pipeline). It's broken into independent per-stage budgets below, so a measured overrun can later be pinned on a specific stage, not just "the system" as a whole. This is a broader budget than the <100ms embedded-feasibility target already set for the sensing stage alone (Table 2). That figure only covers sensing and position-estimation; this one covers everything up to the point this project's own deliverable ends. Sensor acquisition reuses the already-recommended module (Section 2.4.1.2, Section 2.5), not an assumed frame rate. The remaining stages, fusion, inference, output, are estimated from typical figures for comparable embedded pipelines, since this project hasn't built and measured them yet.

Component	Budget	Notes
Sensor acquisition (mmWave radar)	50-100 ms	10-20Hz base configuration, consistent with the "typical demo" figures in Section 2.4.1.2 and with TI's own indoor people-tracking reference design for this chip ([tidue71d]), which achieves ~18Hz (55ms); the mmWave SDK allows configuring higher rates, see Sensor Sampling Rate below.
Sensor fusion (Kalman/EKF over multiple frames)	0-20 ms	Lightweight relative to inference; can run in parallel with the next frame's capture.
Inference (height-category classification)	1-5 ms	Theoretical estimate for a simple geometric approach: take the highest and lowest point in the person's radar point cloud, compare against an adult/child threshold. No trained model. Not yet measured on target hardware, and adult/child classification itself isn't a confirmed VH requirement yet (see below and Non-Camera Person Position & Height Sensing Techniques).
Output exposed to the pipeline (only if sensing and the pipeline consumer run on separate machines)	0 ms local / 50-100 ms networked	Negligible on one local machine; rises materially over Wi-Fi/Ethernet if split across machines, per Non-Camera Person Position & Height Sensing Techniques.
Total, this project's own scope	51-125 ms	Sum of the ranges above, local only. This is the part this project actually builds and can bench-test.

Downstream of that, MindLabs' existing Holobox / Unreal Engine system renders the Virtual Human's actual reaction, one animation frame updating skeletal targets, typically 16-33ms at 30-60 FPS. That's not something this project builds, controls, or can bench-test; Finished Products lists only the embedded sensing system as the prototype deliverable. It's mentioned here purely as context: added to this project's own 51-125ms, it lands around 67-158ms, which is what the <150ms full-pipeline target (Perceptual Latency Thresholds) is actually measuring against.

3.4.2. Jitter Margin

Adding headroom for scheduling jitter and frame misalignment, a planning margin pending the bench measurement in Requirements and Verification, not a measured figure:

Base budget, this project's own scope (mid-point of each range above):	~88 ms
Jitter margin:	+20 ms
Resulting estimate:	~108 ms

~108 ms is the planning target for NFR-Latency-01, this project's own scope, sensing through to exposing the estimate. That leaves roughly 42ms of the full 150ms budget for the downstream Virtual Human rendering this project doesn't build, close to the 16-33ms typically estimated for one render frame. That margin rests on an unvalidated assumption, though. If the simple top/bottom-point approach isn't accurate enough on real, noisy point-cloud data, and a heavier model is needed after all, inference could grow back toward 30-50ms and eat most of that margin. The bench measurement in Requirements and Verification is where this actually gets tested, not assumed.

3.4.3. Sensor Sampling Rate

TI doesn't publish a fixed maximum frame rate for the IWR6843AOP. The mmWave SDK treats frame periodicity as an open, configurable float value ([[mmwavesdkug](#)]). The only real constraints: active-chirp duty cycle stays $\leq 50\%$ of the frame period, and there's enough time for the UART transfer. No documented ceiling. So this chapter plans around a 10-20Hz base rate for the already-selected IWR6843AOPEVM ([Section 2.5](#)), the same "typical demo" range from [Section 2.4.1.2](#), rather than picking one number out of thin air. TI's own indoor people-tracking reference design for this chip ([[tidue71d](#)]) backs that up: a 288-chirp, 1.78GHz-sweep profile, tuned for 8.4cm range resolution over 6m, runs at 55ms (~18Hz) with a 46% duty cycle, within the SDK's $\leq 50\%$ limit. That's a real, published example, not just a demo default. Since frame periodicity is open, running faster is possible and would ease the budget further. But TI hasn't published a faster worked example for this use case. A leaner chirp profile, trading range/velocity resolution for speed, is worth exploring, but that's for the bench measurement in Requirements and Verification to check, not something to assume here.

3.4.4. Compute Platform Implications

Height-category classification has two options: a simple top/bottom-point gap check (no trained model, ~1-5ms), or the keypoint-extraction model from Non-Camera Person Position & Height Sensing Techniques (ProbRadarM3F, [2024](#), ~30-50ms). Neither figure is measured or literature-backed for this task, but ML is likely overkill for a gap this simple. The simple approach is primary; the keypoint model is the fallback if it isn't accurate enough against real, noisy point clouds (occlusion, missed points, clustered people), untested either way. Even the fallback's speed is unproven here: Huang et al. ([2021](#)) show tracking on a Raspberry Pi 4, not classification.

This also assumes the VH's behaviour depends on adult/child category at all, not confirmed yet (Non-Camera Person Position & Height Sensing Techniques). And if inference runs on a paired host, that host also runs the microphone array's beam-steering (Candidate Sensor Technologies, Microphone Array), a compute-contention risk the bench measurement should check for.

3.5. Requirements and Verification

3.5.1. Non-Functional Requirements

NFR-Latency-01	This project's own pipeline (radar frame arrival to output exposed to the Unreal Engine pipeline) shall take < 110ms (p50) and < 150ms (p95), nominal local-processing conditions.
Rationale	The full-pipeline target is <150ms (p50), sensing to visible VH reaction. But VH rendering is MindLabs' existing system, not this project's deliverable (Finished Products). 110ms is this project's own slice, leaving ~40ms for that downstream rendering.
Measurement	Timestamp radar frame arrival, fusion output, inference result, output exposed. Report median/p95 delta, 10-minute window, $N \geq 20$ trials.
Dependencies	Radar frame rate in the 10-20Hz base range (Section 2.4.1.2, [tidue71d]). Inference ≤ 5 ms (simple approach) or ≤ 50 ms (ML fallback); both theoretical. No competing background process during measurement.

3.5.2. Verification Strategy

This NFR isn't confirmed yet; it needs the height-classification pipeline to exist first. Once it does:

1. Inject a test stimulus at a known time (e.g. a person of known height stepping into the tracked zone on cue).
2. Log timestamps: radar frame capture, fusion output, inference result, output exposed.
3. Compute the delta from stimulus to that exposed output, per trial (not to the VH's own reaction, that's out of scope).
4. Repeat for $N \geq 20$ trials over 10 minutes, to capture run-to-run variance, not one lucky run.
5. If inference runs on a paired host, repeat with the microphone array's beam-steering workload running concurrently, to check for compute contention.
6. Report median/p95 latency against NFR-Latency-01.

Local-only first (no network hop), to isolate this from Unreal Engine transmission latency. Real-user field testing and sprint-review feedback, per the Project Plan's research methods, come after, once local measurement confirms the pipeline is in-budget.

3.6. Conclusion

RQ8 sets a full-pipeline target of $< 150\text{ms}$ (p50) / $< 200\text{ms}$ (p95), sensing to visible VH reaction, from general HCI thresholds and ITU-T G.114. But that full reaction isn't this project's deliverable, only the embedded sensing system is (Finished Products); VH rendering is MindLabs' existing Holobox/Unreal Engine system. So NFR-Latency-01 covers a narrower, measurable slice: sensing through to exposing the estimate, $< 110\text{ms}$ (p50) / $< 150\text{ms}$ (p95).

That narrower budget has some room, on paper, under an unvalidated assumption: no documented max frame rate for the IWR6843AOPEVM, so this chapter plans around TI's 10-20Hz base rate (Section 2.4.1.2, [tidue71d], $\sim 18\text{Hz}$), and assumes height classification is a simple top/bottom-point calculation, not a trained model. That leaves $\sim 108\text{ms}$ against the 110ms ceiling, not much room. That margin depends on the simple approach holding up, or falling back to the heavier keypoint model, itself unproven for this task (Huang et al., 2021, show tracking on a Raspberry Pi 4, not classification). The bench measurement in Requirements and Verification is where this gets confirmed.

As with Non-Camera Person Position & Height Sensing Techniques, this chapter is a target and a verification plan, not a finished result. Every figure, frame rate, inference cost, fusion time, margin, is theoretical, pending validation once a working pipeline exists.

4. Embedded Compute Platform Selection

4.1. Introduction

This chapter addresses RQ2 and RQ3 from the Project Plan (Chapter 3.6, Research Questions): implementing the sensor on an embedded platform for robust, low-latency position estimates (RQ2), and deploying a lightweight classification model on that same platform (RQ3).

Non-Camera Person Position & Height Sensing Techniques and its Module Selection chapter already fix the sensor: the TI IWR6843AOPEVM. That EVM's own onboard DSP and Cortex-R4F run TI's people-tracking firmware and produce a tracked point-cloud stream; End-to-End Latency Requirements' Compute Platform Implications section already relies on a host board existing downstream to run height classification and expose the result to the Holobox pipeline, without naming or costing one. Candidate Sensor Technologies excludes that host board from its own cost comparison for the same reason: it is a shared requirement across every sensor candidate, not something that distinguishes one sensor choice from another.

This raises an obvious question: if the EVM already tracks people on-chip, why does this project need a separate board at all? The EVM's onboard DSP and Cortex-R4F could, in principle, also run the classification step itself; TI's mmWave SDK is user-programmable, not fixed-function firmware. What it cannot do is talk to the Holobox pipeline: the EVM's only interfaces are UART-over-USB and a GPIO/breakout header, with no Ethernet or WiFi of its own. Bridging that UART output onto the network is the one part of this pipeline the sensor genuinely cannot do standalone, and it is the reason a host board exists in this design at all.

That bridging role could in principle be filled by the machine already running the Holobox's Unreal Engine pipeline, plugging the EVM into it directly and running classification there as software; NFR-Latency-01 already prefers that local case (0ms local vs. 50-100ms networked, Latency Budget). This chapter does not adopt that option: the Project Plan describes the deliverable as an embedded sensing system with its own hardware and software prototype (Assignment, Finished Products), not a service added to MindLabs' existing show PC, so the board this chapter picks is a discrete embedded unit rather than a re-use of that machine.

This chapter closes that gap: it picks the board that receives the EVM's point-cloud stream, runs the classification stage, and exposes the resulting position/height estimate over the network, and checks it against the latency budget and cost constraints already established. It does not cover the microphone array hardware itself, only whether the chosen board has spare compute for the beam-steering workload once that is added (Candidate Sensor Technologies, Microphone Array). It also does not train or validate the classification model itself; that is RQ3's model-accuracy question, addressed separately once a working pipeline exists.

4.2. Research Approach

This chapter follows the Lab and Library strategies the Project Plan assigns to RQ2 and RQ3 (Chapter 3.7, Research Methods): literature review, available product analysis, and prototyping/benchmarking; the in-situ and model-training parts of those rows come later, once a board is chosen and a pipeline exists to benchmark.

4.2.1. Library: Literature Study

The search here is narrower than Non-Camera Person Position & Height Sensing Techniques: rather than surveying sensing modalities, it looks for published or documented cases that pair this project's already-selected radar family (the TI IWR6843 series) with a specific class of host board, to check whether the interface and compute split this project plans to use has actually been demonstrated, not just assumed.

4.2.2. Lab: Available Product Analysis

Candidate boards are narrowed to specific, currently sourceable hardware, checked against live distributor listings rather than assumed pricing, for the same reason as Non-Camera Person Position & Height Sensing Techniques: an unverified cost estimate is not a usable basis for the budget-approval process in the Project Plan (Chapter 6.1, Finance and Risks). Literature precedent and product availability are again treated as separate questions: a board can be well-documented for general embedded use without a published case of it running this project's specific sensor pipeline.

4.2.3. Lab: Bench Experimentation Strategy

The board recommended here is not bench-tested yet; that happens once the radar module itself is on hand, alongside the phased implementation plan already set out in Non-Camera Person Position & Height Sensing Techniques (position, then height, then multi-person, then speaker detection). The specific risk this chapter flags for that bench stage is the reported USB bandwidth limit on the host interface (Candidate Embedded Platforms, Raspberry Pi 4 Model B), which needs checking against this project's own point-cloud data rate before it can be treated as settled.

4.3. Evaluation Criteria

As with Non-Camera Person Position & Height Sensing Techniques, candidate boards are compared against explicit criteria traced back to a research question, an already-established budget, or a project constraint, rather than a general sense of which board is "more powerful."

Table 6. Evaluation Criteria

Criterion	Description
Host interface compatibility	Must receive the already-selected IWR6843AOPEVM's tracked point-cloud output without added interface hardware, since the sensor module is fixed (Module Selection).
Real-time classification headroom	Must run height classification inside the ~108ms per-project latency budget (Latency Budget), for both the primary simple threshold approach (~1-5ms) and, if it becomes necessary, the ProbRadarM3F keypoint fallback (~30-50ms), per RQ2's low-latency requirement.
Concurrent workload headroom	Must have spare compute if there is necessity to add a microphone array.
Output path to the Holobox pipeline	Needs a network interface so exposing the estimate does not force it onto the same machine as Unreal Engine; local-only adds 0ms, networked adds 50-100ms (NFR-Latency-01).
Literature and community precedent	Must show a documented case pairing this exact radar family with the board class, not just general embedded capability, for the same reason as Evaluation Criteria in Non-Camera Person Position & Height Sensing Techniques.
Hardware cost	Cost of the board, checked against sourceable listings; still a real constraint, since no project budget has yet been agreed with MindLabs (Finance and Risks).
Power and footprint	Must suit mounting near or inside the Holobox installation: passive or small-fan cooling, not a rack-mounted or actively loud unit.
Availability at ISSD	Whether the board can be borrowed directly from Fontys' own ISSD hardware loan desk; since MindLabs is itself a Fontys initiative, a board available there can be obtained without waiting on the mentor budget approval that every purchased candidate still needs (Finance and Risks).

Host interface compatibility and real-time classification headroom are treated as closer to hard requirements: a board that cannot receive the sensor's own output, or cannot fit inside the latency budget already committed to in Latency Budget, is not usable regardless of cost or literature backing. Availability at ISSD is also treated as close to decisive rather than just another weighed factor: it is the only criterion that can remove the pending-budget blocker entirely, rather than just keeping the ask small. The remaining criteria are weighed against each other in the chapter that follows.

4.4. Candidate Embedded Platforms

The IWR6843AOPEVM's own onboard DSP and Cortex-R4F already run TI's people-tracking firmware and produce a tracked point-cloud stream ([tidue71d], [mmwavesdkugl]); the host board this chapter picks does not run the mmWave SDK itself. Its job is narrower: consume that stream, run the classification stage, and expose the result to the Holobox pipeline. Barral et al. ([barral2024mmwaveiot]) document exactly this split on the same radar family: TI's onboard firmware handles point-cloud generation and a first-pass extended-Kalman-filter track, while their paired host runs the downstream filtering, clustering, and output stack.

4.4.1. Bare-Metal MCU

Finance and Risks' own risk table already names an MCU-to-Linux-SBC fallback direction, implying the project could in principle start at the MCU tier. A representative part, the ESP32-S3, is cheap (a few euros) and has built-in WiFi, but no gathered source pairs an MCU-class part with this radar family for point-cloud consumption, and running the sensor's output parser, a classification stage, and later the microphone array's beam-steering concurrently is a level of scheduling and memory management an MCU without a real OS is not well suited to for this project's timeline. This tier is ruled out as a starting point rather than attempted first and abandoned; the project starts directly at the Linux single-board-computer (SBC) tier the risk table already anticipates.

4.4.2. Raspberry Pi 4 Model B (4GB/8GB)

A quad-core Cortex-A72 @ 1.5GHz, Gigabit Ethernet, dual-band WiFi, and USB 3.0/2.0, sold in 2GB (€60.49), 4GB (€108.89), and 8GB (€157.29) RAM variants at current **Kiwi Electronics** (NL) listings, all sharing the same CPU and interfaces.

The 4GB variant is the one carried forward. The primary design does not need a ROS-based stack at all (Barral et al.'s ROS overhead was that paper's own architecture, not a requirement of this one), and classification itself is expected to run in the IWR's own onboard firmware rather than on the host, per Introduction; the host's own job reduces to receiving a small already-classified result over UART, running the microphone array's beam-steering, and bridging both to the network. None of that approaches a 4GB ceiling, so there is no overhead case for paying for 8GB up front. Neither literature source below reports which RAM variant it used, so precedent does not favour one variant over another either way. The one scenario that would actually need more memory, the raw point cloud plus a PyTorch runtime for the ProbRadarM3F fallback, is also a CPU-throughput problem, not just a memory one; if it is ever triggered, escalating to the Raspberry Pi 5 (Recommended Approach) addresses both at once, rather than buying 8GB of RAM on this board now against a scenario that, if it happens, calls for a faster CPU anyway.

This is the only candidate with a literature precedent for the exact combination this project plans to use: Barral et al. ([barral2024mmwaveiot]) pair an IWR6843ISK, the same radar family as the already-selected IWR6843AOPEVM, with a Raspberry Pi 4 over USB, running TI's onboard tracking firmware on the radar and a downstream ROS-based clustering/tracking/output stack on the Pi. Huang et al. ([huang2021mmwave]), already cited in Candidate Sensor Technologies,

separately show real-time multi-person mmWave tracking running on the same board.

Barral et al. also report a concrete open risk worth carrying into the bench-testing stage: they note bandwidth constraints extracting measurements over the USB interface at higher data rates. This project's own point-cloud rate has not been measured yet, so whether that constraint applies here is unconfirmed, not ruled out.

Pros

- Only candidate with a literature precedent for this exact radar-family-plus-host combination (Barral et al.).
- Separate real-time multi-person tracking precedent on the same board (Huang et al.).
- Low absolute cost at €108.89 incl. VAT for the 4GB variant, keeping the initial hardware ask small given no project budget is approved yet (Finance and Risks).
- Standard GPIO/UART-over-USB interface and a mainstream Linux (Debian-based) ecosystem, satisfying the host-interface-compatibility criterion without added interface hardware.
- 4GB is comfortably enough for the host's own reduced job (receiving a small classified result, running the microphone array, bridging to the network), since classification itself runs on the IWR rather than the host.
- The only candidate reviewed that is available to borrow from Fontys' own ISSD hardware loan desk, sidestepping the pending mentor budget approval entirely (Finance and Risks).

Cons

- No onboard GPU/NPU acceleration; ProbRadarM3F can still run on CPU, but whether it stays inside the latency budget on this board is unproven, and it would need the board to fall back to ingesting the raw point cloud rather than the lean classified stream.
- Barral et al.'s reported USB bandwidth constraint is an open risk, unconfirmed against this project's own data rate, and would only actually matter if that same fallback is triggered.

4.4.3. Raspberry Pi 5

A quad-core, out-of-order Cortex-A76 @ 2.4GHz with 8GB RAM, Gigabit Ethernet with PoE+ support, WiFi, USB 3.0, PCIe 2.0, and a 40-pin GPIO/UART header, €192.38 incl. VAT at current [Kiwi Electronics](#) (NL) listings.

No gathered source pairs this specific board with the IWR6843 family; it shares the same GPIO/UART interface and Debian-based Linux stack as the wider Raspberry Pi family, so integrating this project's sensor-parsing and classification software needs no new interface work. Its out-of-order CPU cores and doubled RAM give real headroom for the ProbRadarM3F keypoint fallback (~30-50ms estimate) and for running the microphone array's beam-steering concurrently, though neither is benchmarked on this exact pipeline yet. PCIe 2.0 and USB 3.0 also leave room for a future accelerator add-on, if CPU-only inference on this board still proves insufficient.

Pros

- Meaningfully more CPU and memory headroom than the primary classification approach needs on its own, a margin against the ProbRadarM3F fallback or concurrent microphone-array beam-steering.
- Same GPIO/UART interface and Linux ecosystem as the wider Raspberry Pi family, so no new interface work is needed.

- PCIe and USB 3.0 leave room for a future accelerator add-on if CPU-only inference proves insufficient.

Cons

- No gathered source pairs this specific board with the IWR6843 family.
- €192.38 incl. VAT is a real expense given no project budget is approved yet (Finance and Risks).
- Not available at Fontys' ISSD; adopting it now would mean waiting on mentor budget approval before it can even be obtained.

4.4.4. NVIDIA Jetson Orin Nano Super Developer Kit

A 6-core Cortex-A78AE @ 1.7GHz paired with a 1024-core Ampere GPU and 32 tensor cores (up to 67 TOPS INT8), 8GB LPDDR5, 7-25W power envelope. No EUR-market listing was found at a major distributor; converting Electrokit's (SE) listed 5,329 SEK incl. VAT gives approximately €465, consistent with Aetherix's (NL) directly listed €468.47, so €465-470 is used as the working estimate.

No gathered source pairs this board, or any Jetson model, with the IWR6843 family or with the ProbRadarM3F keypoint model specifically. The case for it is inferred, not demonstrated: ProbRadarM3F ([\[probradarm3f2024\]](#)) is a PyTorch/CUDA model, and Jetson's CUDA and TensorRT stack is the standard embedded deployment target for that model class in general. That is a reasonable inference from how such models are normally deployed, not a result specific to this project's use case.

Pros

- Only candidate with GPU/tensor-core acceleration, the class of hardware the ProbRadarM3F keypoint fallback would most plausibly need if the simple threshold approach underperforms.
- Same GPIO-class headers and USB/networking interfaces as the SBC candidates, so the sensor interface itself is no harder to integrate.

Cons

- No literature or documented case pairs it with this project's specific sensor or classification model.
- At ~€465-470 incl. VAT, this is a substantial per-unit cost against Finance and Risks' unapproved budget.
- A 7-25W power envelope and active cooling requirement, a consideration if the board needs to sit inside or close to the Holobox housing.
- Justified only if the simple threshold classification approach turns out to be insufficient; committing to it now would be pre-paying for a capability not yet shown to be needed.
- Not available at Fontys' ISSD; adopting it now would mean waiting on mentor budget approval before it can even be obtained.

4.4.5. Comparison Summary

Table 7. Candidate Embedded Platform Comparison Summary

Board	Compute Class	Approx . Cost (incl. VAT)	Sensor-Family Precedent	Power Envelope	Available at ISSD
Raspberry Pi 4 Model B (4GB/8GB)	Quad-core Cortex-A72 @ 1.5GHz, CPU-only	€108.89/ €149.99	Yes, direct (Barral et al.; Huang et al.)	~5-8W	Yes
Raspberry Pi 5 (4GB/8GB)	Quad-core Cortex-A76 @ 2.4GHz, CPU-only	€158.07/ €192.38	No (same ecosystem as Pi 4)	~8-12W	No
Jetson Orin Nano Super Dev Kit	6-core Cortex-A78AE + 1024-core Ampere GPU, 67 TOPS INT8	€539	None found	7-25W	No

Figures are board-only, current as of 08-09-2026, and may vary when checked on a different date. Compute-class figures are manufacturer specifications, not measured performance on this project's own pipeline. The Raspberry Pi 4 row uses the 4GB variant's price throughout, the variant recommended below (Raspberry Pi 4 Model B (4GB/8GB)); the 8GB step is available but not needed, since classification runs on the IWR rather than the host and the host's own job stays well inside 4GB. Cost is where the candidates spread out most: the Raspberry Pi 4 (4GB) is the cheapest board reviewed; the Raspberry Pi 5 (8GB) costs roughly 77% more than the Raspberry Pi 4 (4GB); and the Jetson Orin Nano Super costs roughly 5x the Raspberry Pi 4 (4GB). Availability at ISSD is not just another line item: since MindLabs is a Fontys initiative, a board listed there can be borrowed through the college rather than purchased, which matters while no project budget is approved yet (Finance and Risks). Of the three candidates, only the Raspberry Pi 4 is available at ISSD; the Raspberry Pi 5 and the Jetson would both have to be bought, so choosing either now means clearing the budget-approval step first, on top of everything else weighed above.

Table 8. Criteria Comparison Matrix

Board	Precedent	Cost	Power Envelope	Latency Headroom	Concurrent Headroom	GPU/ML Accel.	Available at ISSD	Decision Role
Raspberry Pi 4 Model B (4GB/8GB)	✓	€108.89	~5-8W	✓ (primary); ML fallback unproven on CPU	Unconfirmed	□ (not needed by primary approach)	Yes	Recommended
Raspberry Pi 5	□	€192.38	~8-12W	✓ (primary); ML fallback ~30-50ms est., unbenchmarked	✓	□ (PCIe headroom for future add-on)	No	Fallback, if bandwidth/concurrency risk materializes
Jetson Orin Nano Super Dev Kit	□	€539	7-25W	✓ (incl. ML fallback)	✓	✓	No	Fallback, only if heavier ML model is needed

Interface compatibility and network output are dropped from this matrix: every candidate here uses the same GPIO/UART-over-USB interface and Ethernet/WiFi networking, so those two criteria do not distinguish between them. GPU/ML acceleration and power envelope are added instead, since those are where the three boards actually diverge: the Jetson is the only candidate with dedicated GPU/tensor-core hardware, at roughly 3-5x the power draw of either Raspberry Pi. Precedent and Cost are moved first and a Decision Role column is added because the remaining three criteria (Latency Headroom, Concurrent Headroom, GPU/ML Accel.) are fallback-tier properties, not requirements of the primary design: the primary classification approach runs on the IWR itself and only needs the host to relay a small result, so the Pi 4's □ on GPU/ML acceleration and "Unconfirmed" on concurrent headroom are not deficiencies against a current need, only against a ProbRadarM3F fallback that is itself unconfirmed to be necessary (Recommended Approach). A flat count of checkmarks would favour the Jetson, but the matrix is not weighted evenly: Precedent and Cost decide the starting platform, and the Pi 5/Jetson's extra headroom and acceleration only matter if bench testing later triggers one of the two

fallback paths below. Available at ISSD is included for the same reason as Precedent and Cost: it is a starting-platform criterion, not a fallback one, and the Pi 4 is the only candidate that clears it, since the Pi 5 and Jetson would both need to be purchased rather than borrowed.

4.4.6. Recommended Approach

The Raspberry Pi 4 Model B, 4GB variant, is recommended as the starting platform. It is the only candidate with a literature precedent for the exact split this project plans to use, a TI IWR6843-family radar handling onboard tracking with a paired host running downstream classification and output (Barral et al.), reinforced by a separate real-time-tracking precedent on the same board (Huang et al.). 4GB is sufficient because the host's own job is narrower than that precedent's architecture: classification runs in the IWR's own onboard firmware, not on the host, so the host only needs to receive a small classified result, run the microphone array's beam-steering, and bridge both to the network (Introduction). The 8GB variant exists and costs little more, but there is no overhead case to pay for it now against a workload that does not approach 4GB in the first place. It is also the cheapest board reviewed, so no cost trade-off is being made against a lower-cost alternative. Finance and Risks reports no project budget has been agreed with MindLabs yet, so starting with the cheapest board that already has a working precedent, rather than paying up front for headroom the current design does not need, keeps the initial ask to the company mentor small. It is also the only candidate available to borrow from Fontys' own ISSD hardware loan desk; since MindLabs is a Fontys initiative, this channel is open to the project, and it lets the board be obtained without waiting on that same budget approval at all. The Raspberry Pi 5 and the Jetson are not stocked at ISSD, so adopting either now would mean a purchase, and therefore mentor budget sign-off, before the board could even be picked up; this reinforces the recommendation independently of the cost and precedent case already made above.

This is a starting-point recommendation, not a final hardware order; as with the sensor choice in Candidate Sensor Technologies, it still needs mentor sign-off before use, even though borrowing it from ISSD means no purchase approval is needed to obtain it (Finance and Risks). Two risks carry forward into the bench-testing stage set out in Research Approach: Barral et al.'s reported USB bandwidth constraint, unconfirmed against this project's own point-cloud data rate, and whether the Pi 4 has spare CPU cycles once the microphone array's beam-steering runs alongside classification (Compute Platform Implications). Neither is a RAM problem: the bandwidth constraint is about the UART/USB link, and the concurrent-workload question is about CPU cycles, not memory capacity, so paying for 8GB now would not address either.

If bench measurement shows either risk is real, the Raspberry Pi 5 is the documented next step: same interface and software ecosystem as the Pi 4, so the fallback does not mean re-validating a different platform, at roughly double the compute headroom (per published Geekbench 6.2 figures: ~2.2-2.4x over the Pi 4) for about a 77% cost increase over the Pi 4 (4GB); unlike the Pi 4, it is not available at ISSD, so escalating to it means clearing mentor budget approval first. The Jetson Orin Nano Super Developer Kit stays a further fallback specifically for the ProbRadarM3F keypoint model, if the primary top/bottom-point classification approach turns out not to be accurate enough (Candidate Sensor Technologies); it is not adopted now because that model's need is itself unconfirmed, it is the most expensive and highest-power candidate reviewed at

roughly 5x the Raspberry Pi 4 (4GB)'s cost, and like the Pi 5 it would need to be purchased rather than borrowed from ISSD.

4.5. Conclusion

Candidate Embedded Platforms recommends the Raspberry Pi 4 Model B, 4GB variant, as the host board that receives the IWR6843AOPEVM's already-classified position/height result, runs the microphone array's beam-steering, and bridges both to the Holobox pipeline. It is the only candidate reviewed with a literature precedent for this exact combination (Barral et al.), it is the lowest-cost option among the candidates with any sensor-family precedent, and it fits the latency budget already committed to in Latency Budget. The 4GB variant is sufficient because classification itself runs in the IWR's own onboard firmware rather than on the host (Introduction); the host's own job stays well inside 4GB, so the 8GB variant is available but not needed. It is also the only candidate reviewed that is available to borrow from Fontys' own ISSD hardware loan desk; since MindLabs is a Fontys initiative, this route is open to the project, and it lets the board be obtained without waiting on the mentor budget approval a purchase would need.

As with the sensor recommendation in Candidate Sensor Technologies, this is a starting point pending mentor sign-off before use (Finance and Risks), not a final hardware order; ISSD availability removes the purchase-approval step for obtaining it, not the sign-off itself. Two things carry forward into bench testing once the radar module is on hand: Barral et al.'s reported USB bandwidth constraint, unconfirmed against this project's own data rate, and whether the board has spare compute once the microphone array's beam-steering runs alongside classification. The Raspberry Pi 5 and the Jetson Orin Nano Super Developer Kit remain documented fallbacks, for general headroom and for the ProbRadarM3F keypoint model respectively, neither adopted now because neither risk they would address is confirmed yet, and neither is available at ISSD, so either fallback would additionally mean a purchase and the budget approval that comes with it.

As with the two preceding chapters, this is a recommendation and a verification plan, not a finished result: every figure here is a manufacturer specification or a literature precedent, not a measurement from this project's own hardware.

5. Bibliography

- ITU-T. (2003). *Recommendation G.114: One-way transmission time*. International Telecommunication Union. <https://www.itu.int/rec/T-REC-G.114>
- Shneiderman, B., & Plaisant, C. (2010). *Designing the User Interface: Strategies for Effective Human-Computer Interaction* (5th ed.). Addison-Wesley.
- Alam, M. A. U., Rahman, M. M., & Widberg, J. Q. (2021). *PALMAR: Towards Adaptive Multi-inhabitant Activity Recognition in Point-Cloud Technology*. <https://arxiv.org/pdf/2106.11902>
- Avargel, Y., & Cohen, I. (2011). *Speech measurements using a laser Doppler vibrometer sensor: Application to speech enhancement*. Joint Workshop on Hands-free Speech Communication and Microphone Arrays (HSCMA), Edinburgh, Scotland. <https://israelcohen.com/wp-content/uploads/2018/05/HSCMA2011.pdf>
- Barral, V., Domínguez-Bolaño, T., Escudero, C. J., & García-Naya, J. A. (2024). *An IoT System for Smart Building Combining Multiple mmWave FMCW Radars Applied to People Counting*. arXiv. <https://arxiv.org/html/2401.17949v1>
- Berghi, D., & Jackson, P. J. B. (2023). *Audio Inputs for Active Speaker Detection and Localization via Microphone Array*. <https://arxiv.org/pdf/2307.14739>
- Bian, S., Liu, M., Zhou, B., Lukowicz, P., & Magno, M. (2024). *Body-Area Capacitive or Electric Field Sensing for Human Activity Recognition and Human-Computer Interaction: A Comprehensive Survey*. arXiv. <https://arxiv.org/pdf/2401.06000>
- Bouazizi, M., Mora, A. L., Feghoul, K., & Ohtsuki, T. (2024). *Activity Detection in Indoor Environments Using Multiple 2D Lidars*. *Sensors*. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11154263/>
- Faulkner, N., Parr, B., Alam, F., Legg, M., & Demidenko, S. (2020). *Device Free Localization with Capacitive Sensing Floor* (CapLoc). IEEE Sensors Applications Symposium (SAS). <https://www.researchgate.net/publication/344620665>, building on the earlier CapFloor system by Braun, Heggen, & Wichert (2011).
- Gatto, A., Bassoli, E., Contigiani, M., Frontoni, E., Mancini, A., & Zingaretti, P. (2015). *Energy Harvesting Smart Floor for Indoor People Localization and Tracking*. ASME International Design Engineering Technical Conferences (IDETC-CIE). <https://www.researchgate.net/publication/314616906>
- Huang, X., Cheena, H., Thomas, A., & Tsoi, J. K. P. (2021). *Indoor Detection and Tracking of People Using mmWave Radar*. *Journal of Sensors*, 2021, Article 6657709. <https://onlinelibrary.wiley.com/doi/10.1155/2021/6657709>
- Ibara, K., Hirakawa, M., & Kanetsuna, K. (2013). *Identifying Individuals' Footsteps Walking on a Floor Sensor Device*. *Active Media Technology (AMT 2013)*, Lecture Notes in Computer Science. <https://www.researchgate.net/publication/295102931>
- Klir, S., Lerch, J., Benkner, S., & Khanh, T. Q. (2025). *Multi-Person Localization Based on a Thermopile Array Sensor with Machine Learning and a Generative Data Model*. *Sensors*, 25(2), 419. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11769272/>

- Li, X., Zhang, D., Lv, Q., Xiong, J., Li, S., Zhang, Y., & Mei, H. (2017). *IndoTrack: Device-Free Indoor Human Tracking with Commodity Wi-Fi*. Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies, 1(3), 1-22.
- Miller, R. B. (1968). *Response time in man-computer conversational transactions*. Proceedings of the AFIPS Fall Joint Computer Conference, 33, 267-277.
- Naser, A., Lotfi, A., & Zhong, J. (2021). *Towards human distance estimation using a thermal sensor array*. <https://pmc.ncbi.nlm.nih.gov/articles/PMC8204928/>
- Ngamakeur, K., Yongchareon, S., Yu, J., & Ur Rehman, S. (2020). *A Survey on Device-free Indoor Localization and Tracking in the Multi-resident Environment*. ACM Computing Surveys, 53(4). <https://dl.acm.org/doi/fullHtml/10.1145/3396302>
- Nielsen, J. (1993). *Usability Engineering*. Academic Press. Response-time limits (0.1 s, 1.0 s, 10 s) also summarised by the Nielsen Norman Group: <https://www.nngroup.com/articles/response-times-3-important-limits/>
- Pegoraro, J., & Rossi, M. (2022). *Human Tracking with mmWave Radars: a Deep Learning Approach with Uncertainty Estimation*. <https://arxiv.org/pdf/2205.03274>
- Segers, L., Van Bavegem, D., De Winne, S., Braeken, A., Touhafi, A., & Steenhaut, K. (2015). *An Ultrasonic Multiple-Access Ranging Core Based on Frequency Shift Keying Towards Indoor Localization*. <https://pmc.ncbi.nlm.nih.gov/articles/PMC4570340/>
- Shen, L.-H., Hsiao, A.-H., Chu, F.-Y., & Feng, K.-T. (2023). *Time-Selective RNN for Device-Free Multi-Room Human Presence Detection Using WiFi CSI*. <https://arxiv.org/pdf/2304.13107>
- Texas Instruments. (2020). *People Tracking and Counting Reference Design Using mmWave Radar (TIDUE71D)*. <https://www.ti.com/lit/ug/tidue71d/tidue71d.pdf>
- Texas Instruments. (2022). *mmWave SDK User's Guide (MMWAVE-SDK, Release 3.6)*. https://dr-download.ti.com/software-development/software-development-kit-sdk/MD-PirUeCYr3X/03.06.00.00-LTS/mmwave_sdk_user_guide.pdf
- Shubha, B., & Shastrimath, V. V. D. (2022). *Real-Time Occupancy Detection System Using Low-Resolution Thermopile Array Sensor for Indoor Environment*. IEEE Access, 10, 130981-130995. <https://www.researchgate.net/publication/366350707>
- Su, Y., Hou, H. C., Lan, H., & Ma, C. Z.-H. (2024). *A High-Fidelity mmWave Radar Dataset for Privacy-Sensitive Human Pose Estimation (the mmFree-Pose dataset)*. Bioengineering. <https://pmc.ncbi.nlm.nih.gov/articles/PMC12383697/>
- Wang, G., Zou, Y., Zhou, Z., Wu, K., & Ni, L. M. (2014). *We Can Hear You with Wi-Fi! (WiHear)*. Proceedings of the 20th Annual International Conference on Mobile Computing and Networking (MobiCom).
- Xu, C., Li, Z., Zhang, H., Rathore, A. S., Li, H., Song, C., Wang, K., & Xu, W. (2019). *WaveEar: Exploring a mmWave-based Noise-resistant Speech Sensing for Voice-User Interface*. Proceedings of the 17th Annual International Conference on Mobile Systems, Applications, and Services (MobiSys), 14-26.
- Zhu, B., He, Z., Xiong, W., Ding, G., Huang, T., & Xiang, W. (2024). *ProbRadarM3F: mmWave Radar based Human Skeletal Pose Estimation with Probability Map Guided Multi-Format*

Feature Fusion. <https://arxiv.org/html/2405.05164v2>